

## Intermediate Level

### Advanced SSRF Attack Techniques

#### Bypassing Common Defenses

##### Blacklist Bypass Techniques:

- **IP Address Obfuscation:**
  - Decimal notation: `http://2130706433/` (127.0.0.1)
  - Octal notation: `http://0177.0.0.1/`
  - Hexadecimal notation: `http://0x7f000001/`
  - IPv6 mapping: `http://[::ffff:127.0.0.1]/`
  - IPv6 compressed: `http://[::1]/`
  - IPv6 with embedded IPv4: `http://[0:0:0:0:ffff:127.0.0.1]/`
  - Dword notation: `http://0x7f.0x0.0x0.0x1/`
  - Zero suppression: `http://127.1/`
  - Octal with padding: `http://0127.0.0.1/`
- **Domain Obfuscation:**
  - Using subdomains: `http://localhost.attacker.com/`
  - Registering domains containing target names
  - Using encoded domains: `http://xn--80a.com/` (Punycode)
  - Using domains that resolve to internal IPs
- **URL Encoding Tricks:**
  - Double encoding: `http://127.0.0.1/` → `http%3A%2F%2F127.0.0.1%2F`
  - Mixed encoding: `http://12%37.0.0.1/`
  - UTF-8 characters: `http://①②⑦.①.①.①/`
  - Unicode normalization: Using visually similar characters
- **Redirection Techniques:**
  - Open redirects: `https://example.com/redirect?url=http://internal-service/`
  - 301/302 redirects to internal URLs
  - Meta refresh redirects
  - JavaScript redirects
  - Short URLs that resolve to internal addresses

##### Protocol Smuggling:

- Scheme mixups: `gopher://127.0.0.1:25/`
- URL protocol handler abuse: `java:Runtime.getRuntime().exec('command')`
- Less common protocols:
  - `dict://` (Dictionary Server Protocol)
  - `sftp://` (SSH File Transfer Protocol)
  - `ldap://` (Lightweight Directory Access Protocol)
  - `tftp://` (Trivial File Transfer Protocol)
  - `rtsp://` (Real Time Streaming Protocol)

#### Cloud-Specific Bypass Techniques:

- Alternative metadata service endpoints:
  - AWS: `http://169.254.169.254/`
  - GCP: `http://metadata.google.internal/`
  - Azure: `http://169.254.169.254/metadata/instance`
  - Digital Ocean: `http://169.254.169.254/metadata/v1/`
- Alternative metadata service paths:
  - AWS: `/latest/meta-data/` vs `/latest/user-data/`
  - GCP: `/computeMetadata/v1/` with Metadata-Flavor header
  - Undocumented metadata endpoints

#### DNS Rebinding Techniques:

- Fast-flux DNS: Rapidly changing IP resolution
- Time-based rebinding: Valid IP initially, internal IP later
- Multiple-Answer rebinding: Returns multiple IPs including internal ones
- Advanced tools like Singularity for automated DNS rebinding
- Rebinding with short TTLs to bypass DNS caching

#### Advanced SSRF Exploitation

##### Internal Service Enumeration:

- Port scanning via SSRF:
  - Timing-based port scanning
  - Error-based detection
  - Response size analysis
  - TCP vs UDP port scanning

- Banner grabbing
- Service Fingerprinting:
  - Analyzing response patterns
  - Identifying web servers, application servers
  - Detecting common admin interfaces
  - Finding application version information

#### Exploiting Specific Services:

- Redis exploitation via SSRF:
  - Gopher protocol for raw TCP
  - Writing to `authorized_keys` for access
  - Configuration manipulation
  - Data extraction
- Memcached exploitation:
  - Data extraction
  - Cache poisoning
  - Revealing application secrets
- Elasticsearch exploitation:
  - Reading sensitive indices
  - Extracting configuration
  - Modifying data if writable
- Docker exploitation:
  - Accessing Docker API
  - Creating privileged containers
  - Executing commands in containers

#### Complex SSRF Chains:

- SSRF to XXE:
  - Accessing internal XML parsers
  - Triggering XXE via internal services
  - Reading local files through XXE
- SSRF to RCE:
  - Exploiting local command injection points

- Accessing vulnerable admin interfaces
- Exploiting unpatched internal services
- Attacking deployment systems
- **SSRF to CSRF:**
  - Triggering CSRF on internal applications
  - Changing configurations
  - Executing privileged operations

#### **Out-of-Band Exploitation:**

- **DNS for data exfiltration:**
  - Encoding data in subdomain names
  - Using DNS resolution for blind SSRF confirmation
  - Creating DNS tunnels for command and control
- **HTTP callbacks for confirmation:**
  - Using unique identifiers in requests
  - Blind SSRF validation via external callbacks
  - Burp Collaborator or similar tools for monitoring
- **Time-based techniques:**
  - Measuring response times to infer results
  - Sleep-based confirmation
  - Resource exhaustion techniques

#### **Intermediate Defense Mechanisms**

#### **Improved Validation Approaches:**

- **Context-aware URL validation:**
  - Parse and validate each URL component
  - Validate against hostname patterns, not just exact matches
  - Domain suffix validation
  - Implementing URL parsing correctly across frameworks
- **Defense-in-depth validation:**
  - Combining multiple validation strategies
  - Multi-stage validation pipeline
  - Pre and post-parsing validation

- Static and dynamic validation

#### **Network-Level Controls:**

- **Microsegmentation:**
  - Strict network boundaries between services
  - Zero-trust network model
  - Application-level firewalls
  - Host-based firewalls
- **Service-specific firewall rules:**
  - Restricting outbound connections
  - Protocol-level filtering
  - Blocking connections to internal services
  - Explicit allow-listing for external services
- **Proxies and Gateways:**
  - Dedicated forward proxies for external requests
  - Gateway services for third-party interactions
  - URL rewriting and sanitization
  - Response filtering

#### **Language and Framework-Specific Protections:**

- **Java:**
  - Using `java.net.URL` whitelist
  - `URLConnection` restrictions
  - Custom protocol handler limitations
  - Java Security Manager configurations
- **Python:**
  - `urllib` safe handling
  - Requests library with proxy settings
  - Custom adapter implementation
  - Protocol restriction
- **Node.js:**
  - HTTP client restrictions
  - URL parsing protections

- Custom Agent implementation
- Protocol whitelist enforcement
- Ruby:
  - Net::HTTP safeguards
  - Open-uri restrictions
  - Faraday middleware protections
  - Custom URL validation
- PHP:
  - stream\_context\_create options
  - curl\_setopt restrictions
  - allow\_url\_include and allow\_url\_fopen settings
  - Custom wrapper restrictions

#### Advanced Application Architecture:

- Mediator services:
  - Dedicated services for external communication
  - Enforcing strict contracts for requests
  - Response sanitization and filtering
  - No direct access to internal network
- URL tokenization:
  - Pre-registering allowed destinations
  - Using signed tokens for approved URLs
  - Time-limited URL tokens
  - Encrypted URL parameters
- Request signing:
  - Cryptographic validation of request destinations
  - Chain of trust for request forwarding
  - HMAC-based URL validation
  - Signed request parameters

#### Intermediate Detection and Response

##### Enhanced Detection Patterns:

- Contextual detection:

- Understanding normal application traffic patterns
- Detecting anomalous request destinations
- Profiling allowed external services
- Baselining typical request patterns
- Behavioral analysis:
  - Monitoring request frequency patterns
  - Connection duration analysis
  - Data volume monitoring
  - Request path analysis

#### WAF and Runtime Protection:

- Custom WAF rules:
  - Detecting SSRF patterns in parameters
  - IP-based blocking for internal destinations
  - Protocol restriction enforcement
  - Behavioral rules for request patterns
- RASP solutions:
  - Runtime detection of SSRF attempts
  - Context-aware blocking
  - Real-time intervention
  - Integration with existing security tools

#### Incident Response for SSRF:

- Containment strategies:
  - Temporarily restricting external requests
  - IP blocking for suspected attack sources
  - Disabling vulnerable functionality
  - Routing suspicious traffic through inspection proxies
- Forensic analysis:
  - Identifying compromised endpoints
  - Tracing request paths through systems
  - Analyzing exposed data
  - Determining attack duration and scope

## Intermediate Implementation Examples

### Improved Validation in Python

```
python
```

```
import re
```

```
import socket
```

```
import ipaddress
```

```
from urllib.parse import urlparse
```

```
from requests import get
```

```
def is_valid_external_url(url):
```

```
    """
```

```
    Comprehensive URL validation for SSRF prevention
```

```
    """
```

```
    try:
```

```
        # Basic URL format validation
```

```
        parsed = urlparse(url)
```

```
        if not all([parsed.scheme, parsed.netloc]):
```

```
            return False
```

```
        # Scheme validation - restrict to HTTP/HTTPS
```

```
        if parsed.scheme not in ['http', 'https']:
```

```
            return False
```

```
        # Extract hostname and check against private IP ranges
```

```
        hostname = parsed.netloc.split(':')[0]
```

```
        # Check for localhost in different formats
```

```
        localhost_patterns = [
```

```
            r'^localhost$',
```

```
            r'^127\.',
```

```
            r'^::1$',
```



```
    r'^0\.0\.0\.0$',  
    r'^\.local$'  
]
```

```
for pattern in localhost_patterns:
```

```
    if re.search(pattern, hostname, re.IGNORECASE):
```

```
        return False
```

```
# Resolve hostname to check for internal IPs
```

```
try:
```

```
    ip_addresses = socket.getaddrinfo(hostname, None)
```

```
    for addr_info in ip_addresses:
```

```
        family, _, _, socket_addr = addr_info
```

```
        ip_str = socket_addr[0]
```

```
# Convert to IPv4Address or IPv6Address object
```

```
ip = ipaddress.ip_address(ip_str)
```

```
# Check if IP is private or loopback
```

```
if ip.is_private or ip.is_loopback or ip.is_link_local:
```

```
    return False
```

```
except socket.gaierror:
```

```
# Failed to resolve - could be invalid hostname
```

```
return False
```

```
# Validate against allowed domains (whitelist)
```

```
allowed_domains = ['api.example.com', 'cdn.example.org']
```

```
if hostname not in allowed_domains and not any(hostname.endswith('.') + domain) for domain  
in allowed_domains):
```

```
    return False
```

```
return True
```

```
except Exception as e:
```

```
    # Any parsing or validation error should fail closed
```

```
    print(f"URL validation error: {e}")
```

```
    return False
```

```
def fetch_url_safely(url):
```

```
    """
```

```
    Safely fetch URL content with proper validation
```

```
    """
```

```
    if not is_valid_external_url(url):
```

```
        return "Invalid URL or internal resource requested"
```

```
    try:
```

```
        # Use a timeout to prevent long-running requests
```

```
        response = get(url, timeout=10,
```

```
            # Don't follow redirects to prevent bypass
```

```
            allow_redirects=False,
```

```
            # Limit response size
```

```
            stream=True)
```

```
        # Check redirect attempts
```

```
        if response.status_code in [301, 302, 303, 307, 308]:
```

```
            return "Redirects are not allowed"
```

```
        # Limit response size
```

```
        content = response.raw.read(10 * 1024 * 1024) # 10MB limit
```

```
    return content
```

except Exception as e:

return f"Error fetching URL: {str(e)}"

### Proxy-Based Architecture (Node.js)

javascript

const express = require('express');

const { URL } = require('url');

const axios = require('axios');

const ipRangeCheck = require('ip-range-check');

const dns = require('dns');

const { promisify } = require('util');

const app = express();

const dnsLookup = promisify(dns.lookup);

*// Whitelist of allowed domains*

const ALLOWED\_DOMAINS = ['api.trusted-service.com', 'cdn.approved-domain.org'];

*// Forbidden IP ranges*

const FORBIDDEN\_RANGES = [

'10.0.0.0/8',

'172.16.0.0/12',

'192.168.0.0/16',

'127.0.0.0/8',

'0.0.0.0/8',

'169.254.0.0/16',

'fc00::/7',

'::1/128'

];

async function validateUrl(urlString) {

try {

*// Parse the URL*

```
const url = new URL(urlString);
```

*// Check scheme*

```
if (url.protocol !== 'http:' && url.protocol !== 'https:') {  
  return { valid: false, reason: 'Only HTTP and HTTPS protocols are allowed' };  
}
```

*// Extract hostname*

```
const hostname = url.hostname;
```

*// Check against whitelist*

```
const isWhitelisted = ALLOWED_DOMAINS.some(domain => {  
  return hostname === domain || hostname.endsWith(`.${domain}`);  
});
```

```
if (!isWhitelisted) {  
  return { valid: false, reason: 'Domain not in whitelist' };  
}
```

*// Resolve IP address*

```
try {  
  const { address } = await dnsLookup(hostname);
```

*// Check if IP is in forbidden range*

```
for (const range of FORBIDDEN_RANGES) {  
  if (ipRangeCheck(address, range)) {  
    return { valid: false, reason: 'IP address in forbidden range' };  
  }  
}  
} catch (dnsError) {
```

```

    return { valid: false, reason: 'Failed to resolve hostname' };
  }

  return { valid: true };
} catch (error) {
  return { valid: false, reason: 'Invalid URL format' };
}
}

// Proxy service endpoint
app.get('/proxy', async (req, res) => {
  const url = req.query.url;

  if (!url) {
    return res.status(400).json({ error: 'URL parameter is required' });
  }

  // Validate URL
  const validation = await validateUrl(url);
  if (!validation.valid) {
    return res.status(403).json({ error: `Invalid URL: ${validation.reason}` });
  }

  try {
    // Make the request with strict timeout and size limits
    const response = await axios({
      method: 'get',
      url: url,
      timeout: 5000,
      maxLength: 5 * 1024 * 1024, // 5MB limit
      maxRedirects: 0, // No redirects allowed
    });
  } catch (error) {
    // Handle axios errors
    if (error.response) {
      // Server responded with a status code
      return res.status(error.response.status).json({ error: error.response.data });
    } else if (error.request) {
      // Request made but no response received
      return res.status(500).json({ error: 'No response received' });
    } else {
      // Something else happened while setting up the request
      return res.status(500).json({ error: 'Something else happened' });
    }
  }

  // Return the response data
  return res.json(response.data);
});

```

```

    responseType: 'stream'
  });

  // Forward response headers
  Object.entries(response.headers).forEach(([key, value]) => {
    // Filter out potentially dangerous headers
    if (!['set-cookie', 'transfer-encoding'].includes(key.toLowerCase())) {
      res.setHeader(key, value);
    }
  });

  // Set appropriate content type
  res.setHeader('Content-Type', response.headers['content-type'] || 'application/octet-stream');

  // Stream the response
  response.data.pipe(res);
} catch (error) {
  const statusCode = error.response ? error.response.status : 500;
  const errorMessage = error.response ? `Error: ${error.response.statusText}` : `Request failed: ${error.message}`;

  res.status(statusCode).json({ error: errorMessage });
}
});

const PORT = process.env.PORT || 3000;
app.listen(PORT, () => {
  console.log(`Secure proxy service running on port ${PORT}`);
});

```

Java URL Validation and Secure Client

java

```
import java.net.InetAddress;

import java.net.MalformedURLException;

import java.net.URL;

import java.net.UnknownHostException;

import java.util.Arrays;

import java.util.List;

import java.util.regex.Pattern;

import java.io.IOException;

import java.util.concurrent.TimeUnit;


import org.apache.http.HttpResponse;

import org.apache.http.client.config.RequestConfig;

import org.apache.http.client.methods.HttpGet;

import org.apache.http.impl.client.CloseableHttpClient;

import org.apache.http.impl.client.HttpClientBuilder;

import org.apache.http.util.EntityUtils;
```

```
public class SecureUrlFetcher {
```

```
    // List of allowed domains
```

```
    private static final List<String> ALLOWED_DOMAINS = Arrays.asList(

        "api.trusted-service.com",

        "cdn.approved-domain.org"

    );
```

```
    // Regex patterns for detecting internal IPs
```

```
    private static final Pattern LOCALHOST_PATTERN = Pattern.compile(

        "^(localhost|127\\.\\.0\\.0|0\\.0\\.0\\.0|\\[::1\\]|\\[0:0:0:0:0:0:1\\]).*$",

        Pattern.CASE_INSENSITIVE

    );
```

```

private static final Pattern PRIVATE_IP_PATTERN = Pattern.compile(
    "^(10\\.\\.\\d+\\.\\.\\d+\\.\\.\\d+|" +
    "172\\.\\.(1[6-9]|2\\d|3[0-1])\\.\\.\\d+\\.\\.\\d+|" +
    "192\\.\\.168\\.\\.\\d+\\.\\.\\d+|" +
    "169\\.\\.254\\.\\.\\d+\\.\\.\\d+)$",
    Pattern.CASE_INSENSITIVE
);

```

```

/**

```

```

 * Comprehensive URL validation for SSRF prevention

```

```

 */

```

```

public static boolean isUrlSafe(String urlString) {

```

```

    try {

```

```

        // Basic URL validation

```

```

        URL url = new URL(urlString);

```

```

        // Protocol validation

```

```

        String protocol = url.getProtocol().toLowerCase();

```

```

        if (!protocol.equals("http") && !protocol.equals("https")) {

```

```

            System.out.println("Invalid protocol: " + protocol);

```

```

            return false;

```

```

        }

```

```

        // Get hostname

```

```

        String hostname = url.getHost();

```

```

        // Check for localhost patterns

```

```

        if (LOCALHOST_PATTERN.matcher(hostname).matches()) {

```

```

            System.out.println("Localhost detected: " + hostname);

```

```

            return false;

```

```

        }

```



```
// Validate port (optional)
```

```
int port = url.getPort();
```

```
if (port != -1 && (port < 80 || port > 10000)) {
```

```
    System.out.println("Port outside allowed range: " + port);
```

```
    return false;
```

```
}
```

```
// Domain whitelist check
```

```
boolean isAllowedDomain = false;
```

```
for (String domain : ALLOWED_DOMAINS) {
```

```
    if (hostname.equals(domain) || hostname.endsWith("." + domain)) {
```

```
        isAllowedDomain = true;
```

```
        break;
```

```
    }
```

```
}
```

```
if (!isAllowedDomain) {
```

```
    System.out.println("Domain not in whitelist: " + hostname);
```

```
    return false;
```

```
}
```

```
// Resolve IP addresses
```

```
try {
```

```
    InetAddress[] addresses = InetAddress.getAllByName(hostname);
```

```
    for (InetAddress address : addresses) {
```

```
        String ip = address.getHostAddress();
```

```
// Check for private IP addresses
```

```
    if (PRIVATE_IP_PATTERN.matcher(ip).matches()) {
```

```
        System.out.println("Private IP detected: " + ip);
```

```

        return false;
    }

    // Check for loopback addresses
    if (address.isLoopbackAddress()) {
        System.out.println("Loopback address detected: " + ip);
        return false;
    }

    // Check for link local addresses
    if (address.isLinkLocalAddress()) {
        System.out.println("Link local address detected: " + ip);
        return false;
    }
}

} catch (UnknownHostException e) {
    System.out.println("Failed to resolve hostname: " + hostname);
    return false;
}

return true;

} catch (MalformedURLException e) {
    System.out.println("Malformed URL: " + e.getMessage());
    return false;
}
}

/**
 * Securely fetch content from a URL with proper safeguards
 */

```

```

public static String fetchUrl(String urlString) {
    if (!isUrlSafe(urlString)) {
        return "URL failed security validation";
    }

    CloseableHttpClient httpClient = null;
    try {
        // Configure request with timeouts and restrictions
        RequestConfig config = RequestConfig.custom()
            .setConnectTimeout(5000)
            .setConnectionRequestTimeout(5000)
            .setSocketTimeout(5000)
            .setRedirectsEnabled(false) // Prevent redirects
            .setCircularRedirectsAllowed(false)
            .build();

        // Create HTTP client with configuration
        httpClient = HttpClientBuilder.create()
            .setDefaultRequestConfig(config)
            .setConnectionTimeToLive(10, TimeUnit.SECONDS)
            .build();

        // Create and execute request
        HttpGet request = new HttpGet(urlString);
        HttpResponse response = httpClient.execute(request);

        // Check for redirects
        int statusCode = response.getStatusLine().getStatusCode();
        if (statusCode >= 300 && statusCode < 400) {
            return "Redirects are not allowed";
        }
    }
}

```

```

        // Limit response size (5MB)

        int maxSize = 5 * 1024 * 1024;

        if (response.getEntity().getLength() > maxSize) {
            return "Response too large";
        }

        // Process response

        String responseBody = EntityUtils.toString(response.getEntity());

        return responseBody;

    } catch (IOException e) {
        return "Error fetching URL: " + e.getMessage();
    } finally {
        if (httpClient != null) {
            try {
                httpClient.close();
            } catch (IOException e) {
                // Ignore close errors
            }
        }
    }
}

```

```

public static void main(String[] args) {
    // Example usage

    String url = "https://api.trusted-service.com/data";

    if (isUrlSafe(url)) {
        String content = fetchUrl(url);

        System.out.println("Content: " + content);
    } else {

```

```
        System.out.println("URL failed security validation");  
    }  
}  
}
```